

Summary: in this tutorial, you will learn about the MySQL window functions and their useful applications in solving analytical query challenges.

MySQL has supported window functions since version 8.0. The window functions allow you to solve query problems in new, easier ways and with better performance.

Suppose that we have the `sales` table which stores the sales by employees and fiscal years:

```
CREATE TABLE sales(
    sales_employee VARCHAR(50) NOT NULL,
    fiscal_year INT NOT NULL,
    sale DECIMAL(14,2) NOT NULL,
    PRIMARY KEY(sales_employee,fiscal_year)
);

INSERT INTO sales(sales_employee,fiscal_year,sale)
VALUES( 'Bob',2016,100),
      ('Bob',2017,150),
      ('Bob',2018,200),
      ('Alice',2016,150),
      ('Alice',2017,100),
      ('Alice',2018,200),
      ('John',2016,200),
      ('John',2017,150),
      ('John',2018,250);

SELECT * FROM sales;
```

	sales_employee	fiscal_year	sale
▶	Alice	2016	150.00
	Alice	2017	100.00
	Alice	2018	200.00
	Bob	2016	100.00
	Bob	2017	150.00
	Bob	2018	200.00
	John	2016	200.00
	John	2017	150.00
	John	2018	250.00

It's probably easier to understand window functions is to start with aggregate functions.

Aggregate functions summarize data from multiple rows into a single result row. For example, the following `SUM()` function returns the total sales of all employees in the recorded years:

```
SELECT
    SUM(sale)
FROM
    sales;
```

	SUM(sale)
▶	1500.00

The `GROUP BY` clause allows you to apply aggregate functions to a subset of rows. For example, you may want to calculate the total sales by fiscal years:

```
SELECT
    fiscal_year,
    SUM(sale)
FROM
    sales
GROUP BY
    fiscal_year;
```

	fiscal_year	SUM(sale)
▶	2016	450.00
	2017	400.00
	2018	650.00

In both examples, the aggregate functions reduce the number of rows returned by the query.

Like the aggregate functions with the `GROUP BY` clause, window functions also operate on a subset of rows but they do not reduce the number of rows returned by the query.

For example, the following query returns the sales for each employee, along with total sales of the employees by fiscal year:

```
SELECT
    fiscal_year,
    sales_employee,
    sale,
    SUM(sale) OVER (PARTITION BY fiscal_year) total_sales
FROM
    sales;
```

	fiscal_year	sales_employee	sale	total_sales
▶	2016	Alice	150.00	450.00
	2016	Bob	100.00	450.00
	2016	John	200.00	450.00
	2017	Alice	100.00	400.00
	2017	Bob	150.00	400.00
	2017	John	150.00	400.00
	2018	Alice	200.00	650.00
	2018	Bob	200.00	650.00
	2018	John	250.00	650.00

In this example, the `SUM()` function works as a window function that operates on a set of rows defined by the contents of the `OVER` clause. A set of rows to which the `SUM()` function applies is referred to as a window.

The `SUM()` window function reports not only the total sales by fiscal year as it does in the query with the `GROUP BY` clause, but also the result in each row, rather than the total number of rows returned.

Note that window functions are performed on the result set after all `JOIN` , `WHERE` , `GROUP BY` , and `HAVING` clauses and before the `ORDER BY` , `LIMIT` and `SELECT DISTINCT` .

Window function syntax

The general syntax of calling a window function is as follows:

```
window_function_name(expression) OVER (  
    [partition_definition]  
    [order_definition]  
    [frame_definition]  
)
```

In this syntax:

- First, specify the window function name followed by an expression.
- Second, specify the `OVER` clause which has three possible elements: partition definition, order definition, and frame definition.

The opening and closing parentheses after the `OVER` clause are mandatory, even with no expression, for example:

```
window_function_name(expression) OVER()
```

partition_clause syntax

The `partition_clause` breaks up the rows into chunks or partitions. Two partitions are separated by a partition boundary.

The window function is performed within partitions and re-initialized when crossing the partition boundary.

The `partition_clause` syntax looks like the following:

```
PARTITION BY <expression> [{,<expression>...}]
```

You can specify one or more expressions in the `PARTITION BY` clause. Multiple expressions are separated by commas.

order_by_clause syntax

The `order_by_clause` has the following syntax:

```
ORDER BY <expression> [ASC|DESC], [{,<expression>...}]
```

The `ORDER BY` clause specifies how the rows are ordered within a partition. It is possible to order data within a partition on multiple keys, each key is specified by an expression. Multiple

expressions are also separated by commas.

Similar to the `PARTITION BY` clause, the `ORDER BY` clause is also supported by all the window functions. However, it only makes sense to use the `ORDER BY` clause for order-sensitive window functions.

frame_clause syntax

A frame is a subset of the current partition. To define the subset, you use the frame clause as follows:

```
frame_unit {<frame_start>|<frame_between>}
```

A frame is defined with respect to the current row, which allows a frame to move within a partition depending on the position of the current row within its partition.

The frame unit specifies the type of relationship between the current row and frame rows. It can be `ROWS` or `RANGE`. The offsets of the current row and frame rows are the row numbers if the frame unit is `ROWS` and row values the frame unit is `RANGE`.

The `frame_start` and `frame_between` define the frame boundary.

The `frame_start` contains one of the following:

- `UNBOUNDED PRECEDING` : frame starts at the first row of the partition.
- `N PRECEDING` : a physical N of rows before the first current row. N can be a literal number or an expression that evaluates to a number.
- `CURRENT ROW` : the row of the current calculation

The `frame_between` is as follows:

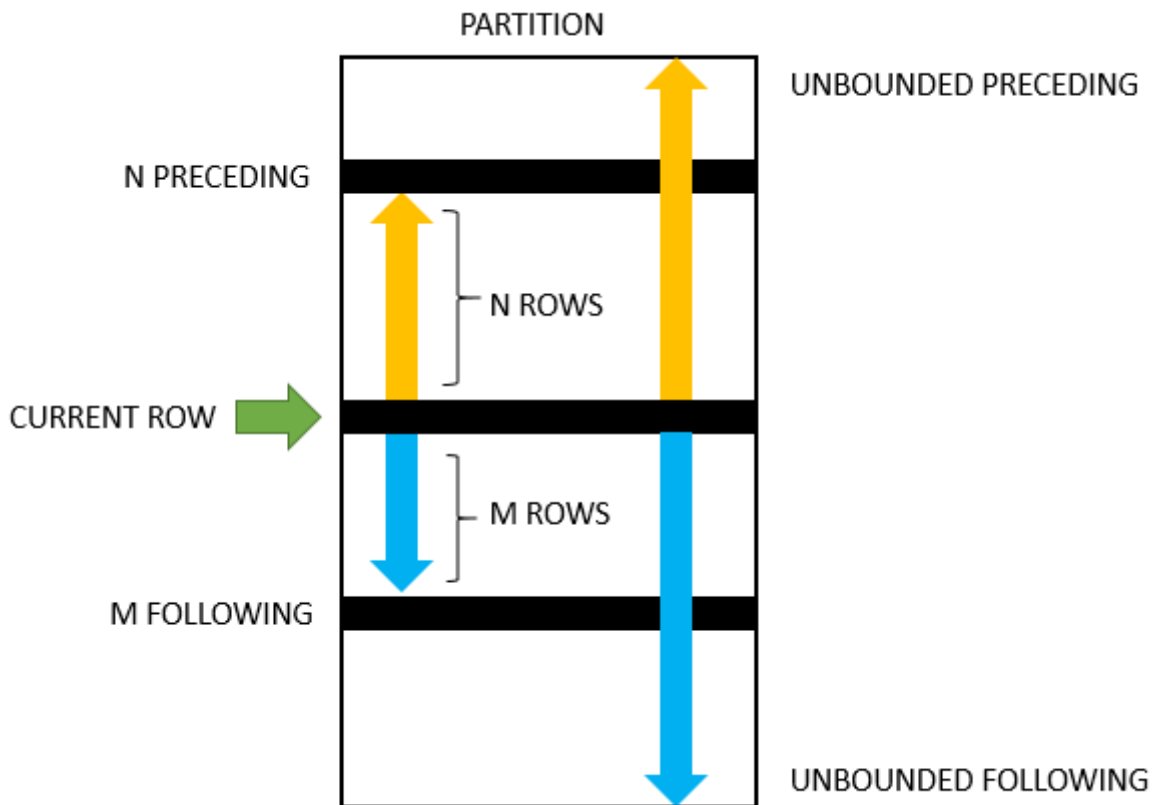
```
BETWEEN frame_boundary_1 AND frame_boundary_2
```

The `frame_boundary_1` and `frame_boundary_2` can each contain one of the following:

- `frame_start` : as mentioned previously.
- `UNBOUNDED FOLLOWING` : the frame ends at the final row in the partition.
- `N FOLLOWING` : a physical N of rows after the current row.

If you don't specify the `frame_definition` in the `OVER` clause, then MySQL uses the following frame by default:

```
RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW
```



MySQL Window Function list

The following table shows the window functions in MySQL:

Type a function name to search...

Name	Description
CUME_DIST	Calculates the cumulative distribution of a value in a set of values.
DENSE_RANK	Assigns a rank to every row within its partition based on the ORDER BY clause. It assigns the same rank to the rows with equal values. If two or more rows have the same rank, then there will be no gaps in the sequence of ranked values.
FIRST_VALUE	Returns the value of the specified expression with respect to the first row in the window frame.
LAG	Returns the value of the Nth row before the current row in a partition. It returns NULL if no preceding row exists.
LAST_VALUE	Returns the value of the specified expression with respect to the last row in the window frame.
LEAD	Returns the value of the Nth row after the current row in a partition. It returns NULL if no subsequent row exists.

Name	Description
NTH_VALUE	Returns value of argument from Nth row of the window frame
NTILE	Distributes the rows for each window partition into a specified number of ranked groups.
PERCENT_RANK	Calculates the percentile rank of a row in a partition or result set
RANK	Similar to the DENSE_RANK () function except that there are gaps in the sequence of ranked values when two or more rows have the same rank.
ROW_NUMBER	Assigns a sequential integer to every row within its partition

In this tutorial, you have learned about the MySQL window functions and their syntax. In the next tutorials, you will learn more about each window function and its applications in more detail.